




Xi'an Jiaotong-Liverpool University

西交利物浦大學

CPT104 - Operating Systems Concepts

Virtual Memory





Virtual Memory

- Background
- Demand Paging
- Copy-on-Write in Operating System
- Page Replacement
- Frame Allocation. Thrashing





Background

Virtual memory (VM) is a method that manages the exceeded size of larger processes as compared to the available space in the memory.

Virtual memory - separation of user logical memory from physical memory.

- ***Only part of the program needs to be in memory for execution.***
- The components of a process that are **present in the memory** are known as **resident** set of the process
- ***Need to allow pages/segments to be swapped in and out.***



Swap space / Swap partition

The implementation of a VM system requires both **hardware** and **software** components.

- The **software** implementing the VM system is known as **VM handler**.
- The **hardware** support is the **memory management unit** built into the CPU.

The VM system realizes a huge memory only due to the hard disk.

With the help of the hard disk, the VM system is able to manage larger-size processes or multiple processes in the memory.

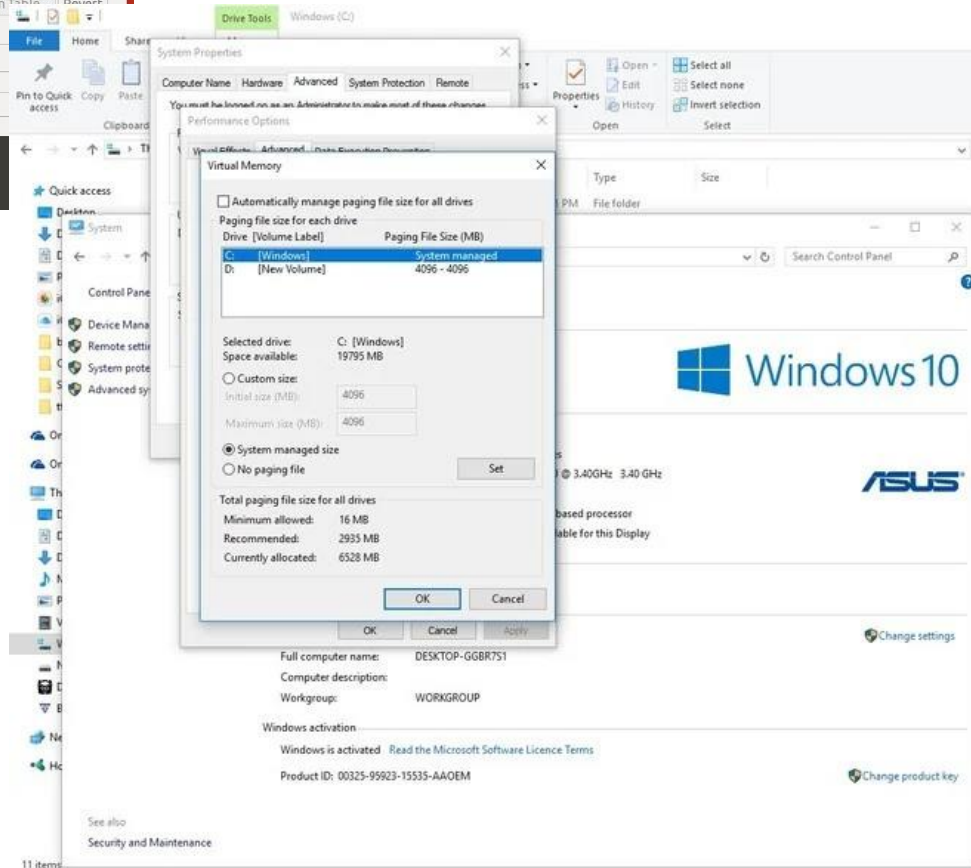
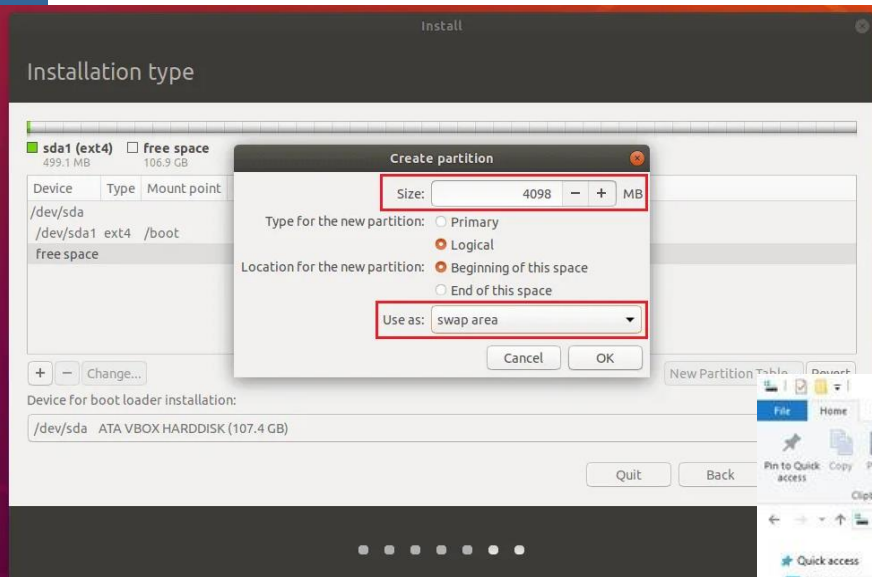
For this purpose, a **separate space** known as **swap space** is reserved in the disk.



Swap space / Swap partition

Create **Swap Partition** – Ubuntu 18.04 LTS

Swap space reserved in the disk.



Windows 10 – **Virtual Memory**



Background


What if there isn't enough physical memory available to load a program?

- A large portion of program's code may be unused
- A large portion of program's code may be used infrequently
- **Idea:** load a component of a process only if it is needed

Virtual memory can be implemented via:

- **Demand paging**

- **Demand segmentation**



Demand loading



Demand Paging

The concept of loading only a part of the process (pages) into memory for processing.






Demand Paging

- when the process begins to run, its pages are brought into memory only as they are needed, and if they're never needed, they're never loaded.
- when a logical address generated by a process points to a page that is not in memory.

□ **Lazy swapper** – never swaps a page into memory **unless page will be needed**

❖ Swapper that deals with pages is a *pager*





Demand Paging

How to recognize whether a page is present in the memory?

- A **bit** set to “**valid**” → the associated page **is** both **valid** (in the logical address space of the process) and **in memory**.
- A **bit** set to “**invalid**” → the page either is **not valid** (not in the logical address space of the process) or **is valid but is currently on the disk**.

What happens if the process tries to access a page that is not in the memory?

- when the page referenced is not present in the memory → **PAGE FAULT**.
- while translating the address through the page table notices that the page-table entry has an invalid bit. **It causes a trap to the OS** so that a page fault can be noticed.

What happens if there is no free frame in the memory ?

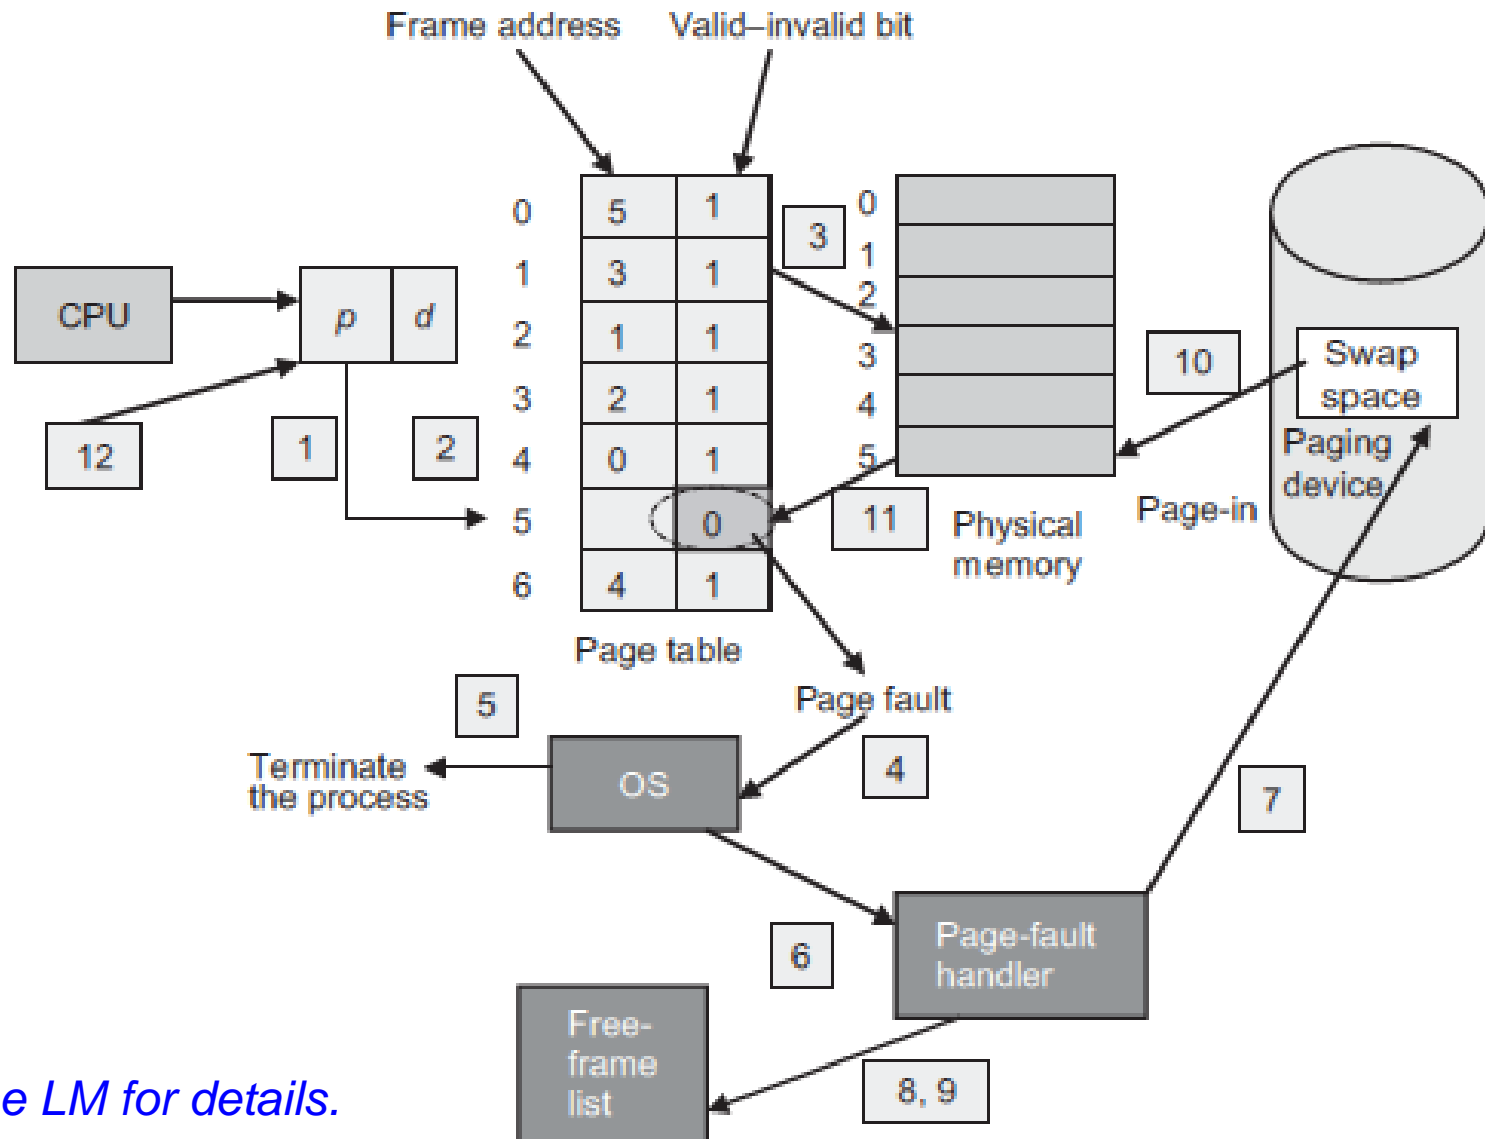
- an existing page in the memory needs to be **paged-out**.
- which page will be replaced? → **page-replacement algorithms**.

valid-invalid bit	
frame	bit
0	4 v
1	i
2	6 v
3	i
4	i
5	9 v
6	i
7	i

page table



Steps in Handling a Page Fault





Performance of Demand Paging

NO page fault: the effective access time = the memory access time.

Otherwise...

p be the probability of a PAGE FAULT $0 \leq p \leq 1$

- if $p = 0$, no page faults
- if $p = 1$, every reference is a fault


The Effective Access Time (EAT):

$$\text{EAT} = (1 - p) \times \text{Memory Access Time} + p \times \text{Page Fault Time}$$

Major components of the page-fault (service) time:

1. Service the page-fault interrupt.
2. Read in the page.
3. Restart the process.

VM system also uses TLB to reduce the memory accesses and increase the system performance.



Copy-on-Write in Operating System

Only pages that are written need to be copied.






Copy-on-Write (COW)

Process creation using the **fork()** system call **may** (*initially*) **bypass the need for demand paging** by using a technique similar to **page sharing**.

Copy-on-Write = strategy that those pages that are never written need not be copied. Only the pages that are written need be copied.

The parent and child process to **share the same pages of the memory** initially.

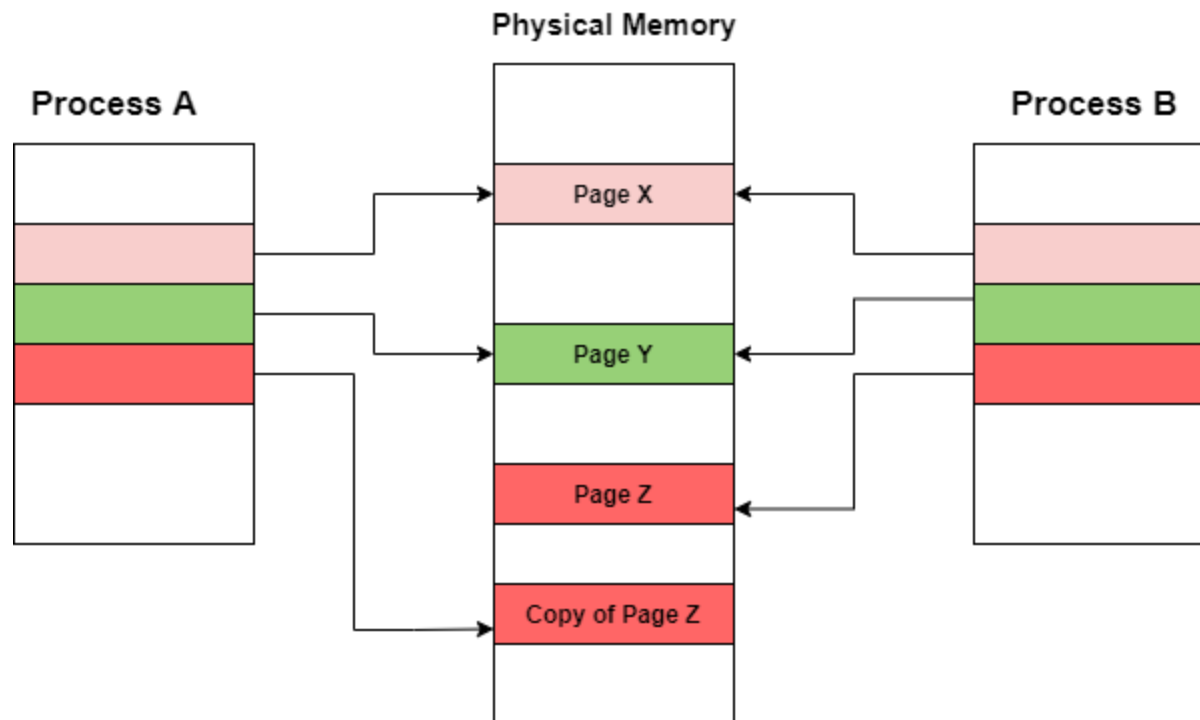


Process A creates a new process - Process B



Copy-on-Write (COW)

If any process either parent or child modifies the shared page, only then the page is copied.



process A wants to modify a page (Z) in the memory



Page Replacement

When a page fault occurs during the execution of a process, a page needs to be paged into the memory from the disk.





Two major problems to implement demand paging:

What happens if there is no free frame? \Rightarrow Page replacement

- When a page is to be replaced, which frame shall be the "*victim*"?
- How to select a replacement algorithm? it wants one with the lowest page-fault rate.

How many frames shall be allocated to each process? \Rightarrow Frame allocation

- When page replacement is required, it must select the frames that are to be replaced.



Page Replacement

The degree of multiprogramming increases \Rightarrow over-allocating memory \Rightarrow NO free frames on the free-frame list, all memory is in use.

A good replacement algorithm achieves:

❑ a *low page fault rate*


- ensure that heavily used pages stay in memory
- the replaced page should not be needed for some time

❑ a *low latency of a page fault*

- efficient code
- replace pages that do not need to be written out
- a special bit called the **modify** (dirty) **bit** can be associated with each page.



Need For Page Replacement






Basic Page Replacement

1. **Find the location of the desired page on disk.**
2. **Find a free frame:**
 - If there is a free frame → use it
 - If there is no free frame → use a *page replacement algorithm* to select a **victim frame**
 - Check the ***modify (dirty) bit*** with each page or frame.
 - If the bit is set** → the page has been modified.
 - If the bit is not set** → the page has not been modified. It need not be paged-out for replacement and can be overwritten by another page because its copy is already on the disk. This mechanism reduces the page-fault service time.
3. **Bring** the desired **page into** the (newly) **free frame** and **update** the **page and frame tables**.
4. **Continue** the process by restarting the instruction that caused the trap.



Page Replacement





Page-replacement algorithms

- ❑ First-In First-Out (FIFO) Algorithm
- ❑ Optimal Algorithm
- ❑ Least Recently Used (LRU) Algorithm
- ❑ Second-Chance (Clock) Algorithm
- ❑ Counting Algorithms

Algorithm evaluation:

- run a **string of memory references** and **compute** the number of **page faults**.
- **recording** a trace of the **pages accessed by a process**.

Reference string - is the sequence of pages being referenced.



First-In First-Out (FIFO) Algorithm

- When a page must be replaced, **the oldest page is chosen**.
- Reference string: **7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1**
- 3 frames** (3 pages can be in memory at a time per process)

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2																		
	0	0	0																		
		1	1																		

page frames

- 15 page-fault



Optimal Algorithm

- Replace page that **will not be used for longest period of time**
 - **9 page-fault**

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		2		2		2						7		
	0	0	0		0		4		0		0						0		
		1	1		3		3		3		1						1		

page frames

Optimal algorithm guarantees the lowest possible page fault rate for a fixed number of frames.

It cannot be implemented - there is no provision in the OS to know the future memory references.

The idea is to predict future references based on the past data,



Least Recently Used (LRU) Algorithm

- replace the page that has been unused for the longest time.

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0		1		1		1
	0	0	0		0		0	0	3	3		3		0		0
		1	1		3		3	2	2	2		2		2		7

page frames

- 12 faults** – better than FIFO but worse than OPT
- Generally good algorithm and frequently used





How to implement LRU replacement?

How to find out a page that has not been used for the longest time?

Two implementations are feasible:

- **Counter** implementation
- **Stack** implementation

COUNTER implementation

- associate with each page-table entry a **time-of-use field** or a **counter**; every time page is referenced through this entry, copy the clock into the counter
- **When a page needs to be changed, look at the counters to find the smallest value**
- **replace the page with the smallest time value.**




How to implement LRU replacement?

STACK implementation

- whenever a page is referenced, it is removed from the stack and put on the top.
- the **most recently used page is always at the top of the stack** and the **least recently used page is always at the bottom**

reference string


4 7 0 7 1 0 1 2 1 2 7 1 2



stack
before
a



stack
after
b





LRU Approximation Algorithms

- LRU needs special hardware and still slow
- **Reference bit** - will say whether **the page has been referred in the last clock cycle or not.**
- **The reference bit for a page is set by the hardware** whenever that **page is referenced** (either a read or a write to any byte in the page).
- **Reference bits are associated with each entry in the page table.**
 - With each page associate a bit, **initially = 0**
 - When **page is referenced**, bit set to 1





Second-Chance (Clock) Page-Replacement Algorithm

- keeps a circular list of pages, with the "**hand**" (**iterator**) pointing to the last examined page in the list.

RB = Reference bit or **Use bit**

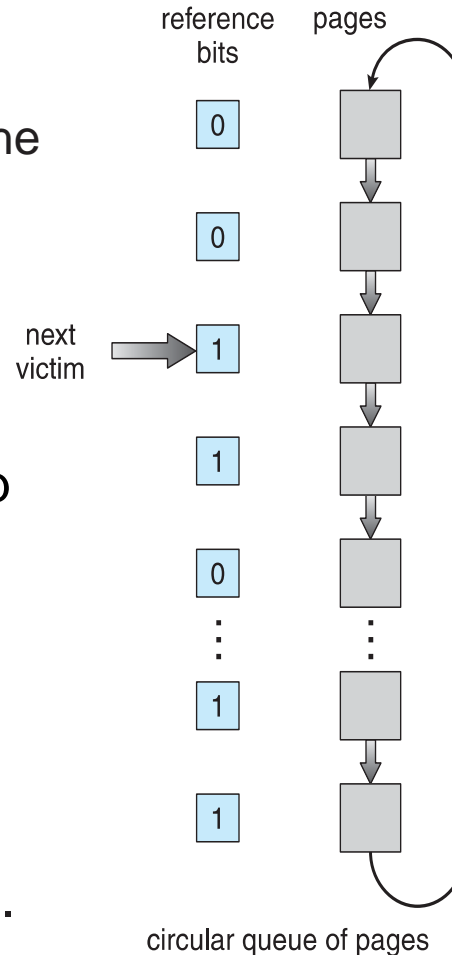
give information regarding whether the page has been used

Iterator scan:


- IF page's RB = 1, set to 0 & skip
- ELSE if RB = 0, remove

The idea behind:

A page that is being frequently used will not be replaced (RB=1).



(a)




(b)



Counting-Based Page Replacement

- Keep a counter of the number of references that have been made to each page.
- **Least Frequently Used (LFU) Algorithm** replaces page with **smallest count**.
- **Most Frequently Used (MFU) Algorithm** is based on the argument that the **page with the smallest count was probably just brought in and has yet to be used**





Frame Allocation. Trashing





Frame Allocation Algorithms

The two algorithms commonly used to allocate frames to a process:

- ❑ **Equal allocation** - In a system with x frames and y processes, each process gets equal number of frames
 - ❑ *Example*, if there are 100 frames (after allocating frames for the OS) and 5 processes, give each process 20 frames
- ❑ **Proportional allocation** – Frames are allocated to each process according to the process size.
 - ❑ Dynamic as degree of multiprogramming, process sizes change

s_i = size of process p_i

$$S = \sum s_i$$

m = total number of frames

$$a_i = \text{allocation for } p_i = \frac{s_i}{S} \times m$$



Frame Allocation Algorithms

Proportional allocation - *example*

A system with 62 frames

□ P1 = 10KB

□ P2 = 127KB

Then

□ P1 will get $(10 / 137) * 62 = 4$ frames

□ P2 will get $(127 / 137) * 62 = 57$ frames.





Thrashing

- If a process does not have “enough” pages, **the page-fault rate is very high.**
- This leads to:
 - ▶ Low CPU utilization
 - ▶ Operating system thinks that it needs to increase the degree of multiprogramming
 - ▶ Another process added to the system
- **Thrashing** \equiv a process is busy swapping pages in and out
 - a process spends more time paging than executing



End of Lecture

■ Summary

- Background
- Demand Paging
- Copy-on-Write in Operating System
- Page Replacement
- Frame Allocation. Thrashing

■ Reading

- Textbook 9th edition, **chapter 9 of the module textbook**

