

Operating Systems Concepts

Review I

Final exam structure

I. Fundamentals

II. CPU scheduling, Memory management, Disk scheduling

III. Resource allocation

IV. Operating System in C Language

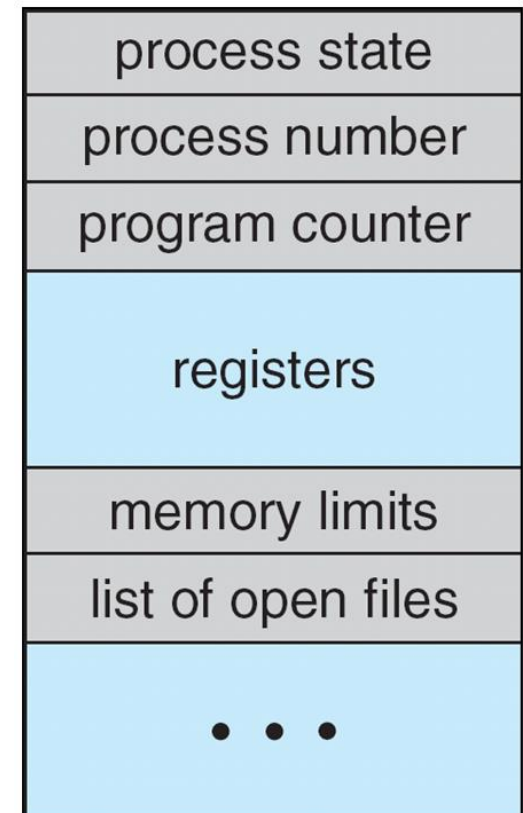
Types of resources managed by an OS

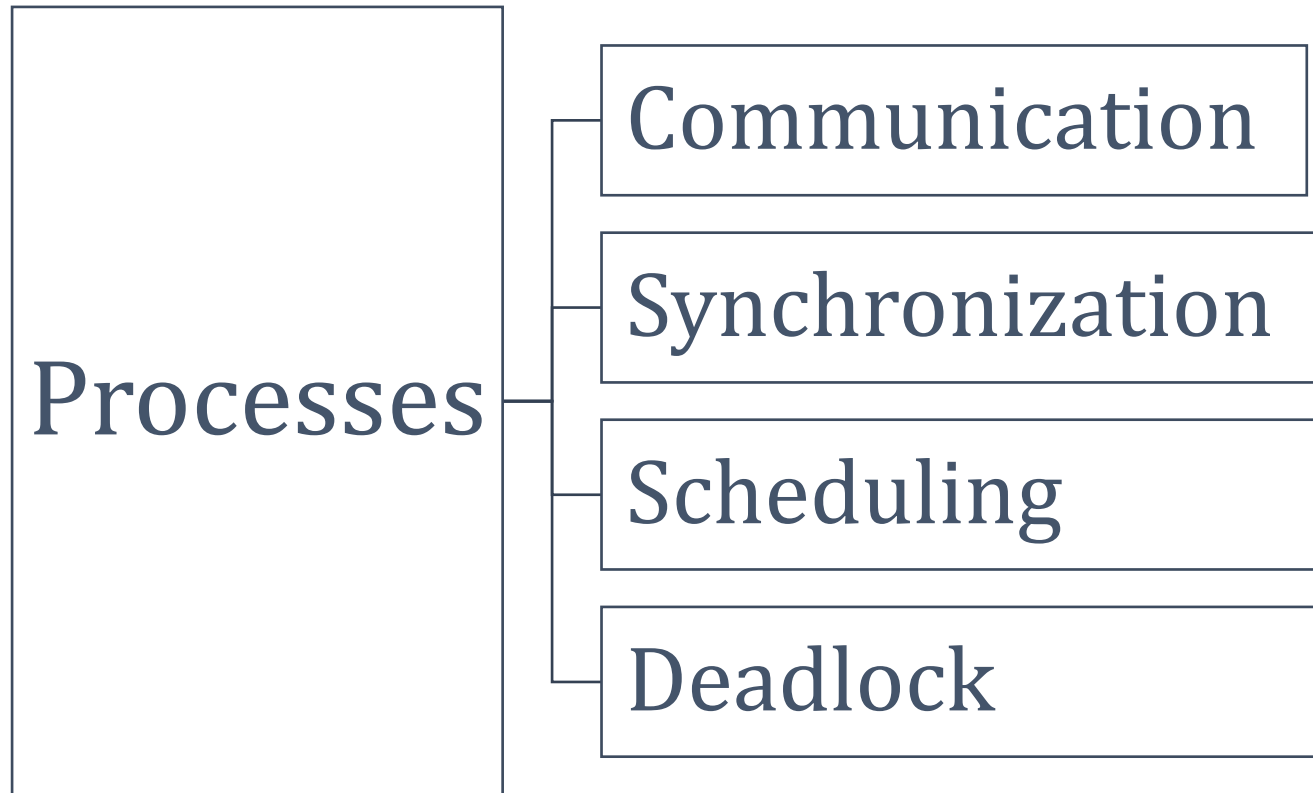
- CPUs (processes)
- main memory and virtual memory
- secondary storage
- I/O devices
- file system and user interface
- communications
- provides protection and security

PROCESSES

- **Process** – a **program in execution**; process execution must progress in sequential fashion
- The state of a process is defined in part by the **current activity of that process**.
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a processor
 - **terminated**: The process has finished execution

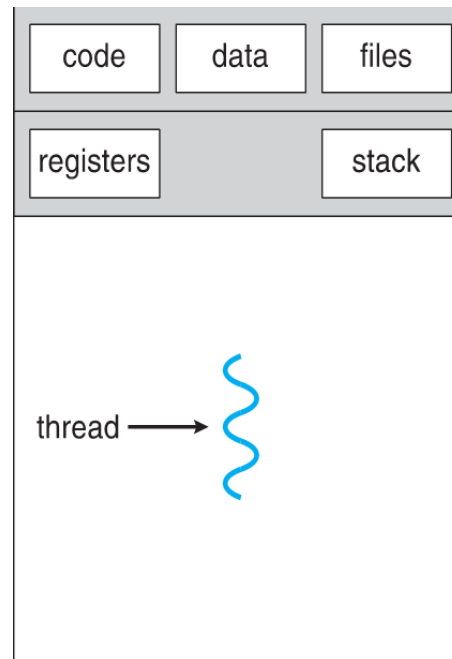
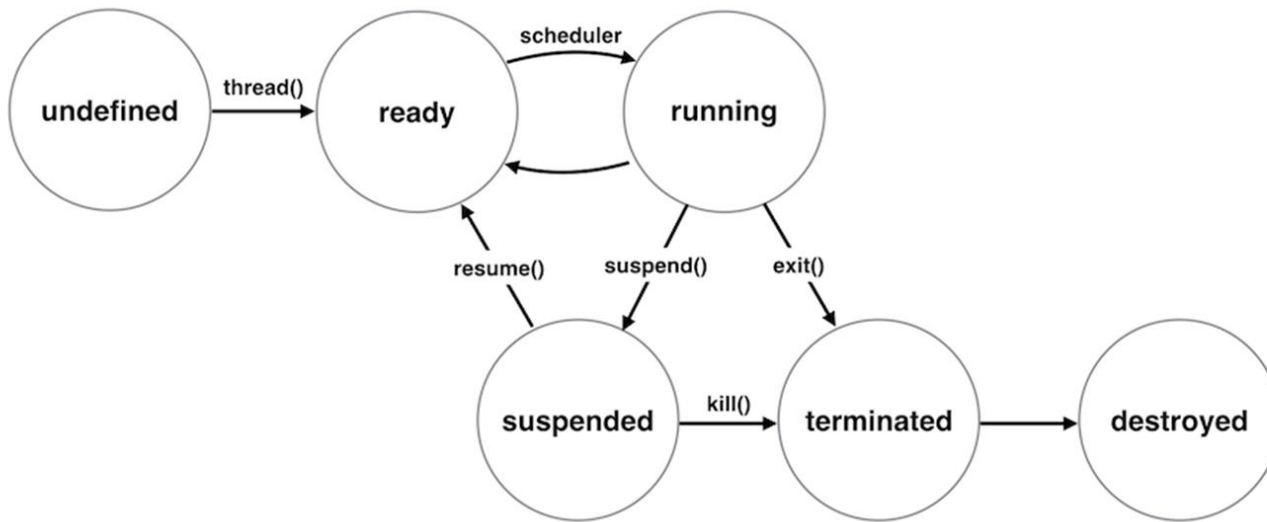
- ❑ **Process Control block** is a data structure used for storing the information about a process.
- ❑ **Each & every process is identified by its own PCB.**
- ❑ It is also called as **context of the process.**
- ❑ PCB of each process resides in the main memory.
- ❑ PCB of all the processes are present in a linked list.



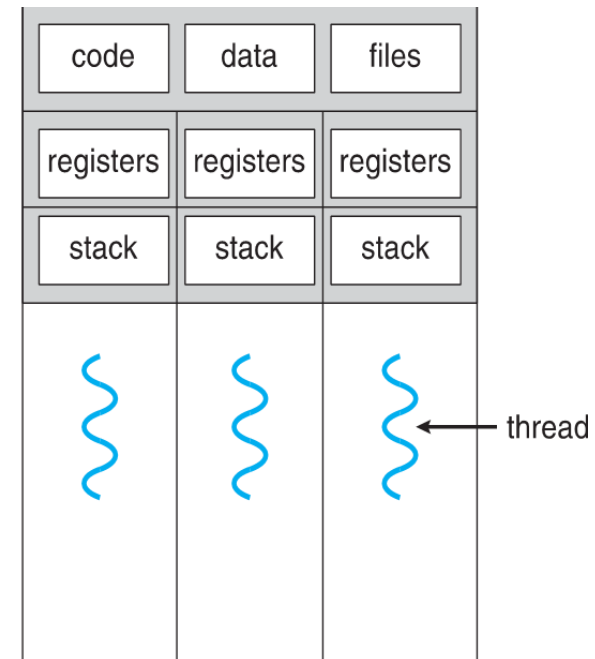


THREADS

- A fundamental unit of CPU utilization that forms the basis of multithreaded computer systems
- Multiple threads can exist within one process, executing concurrently and sharing resources such as memory, while different processes do not share these resources.

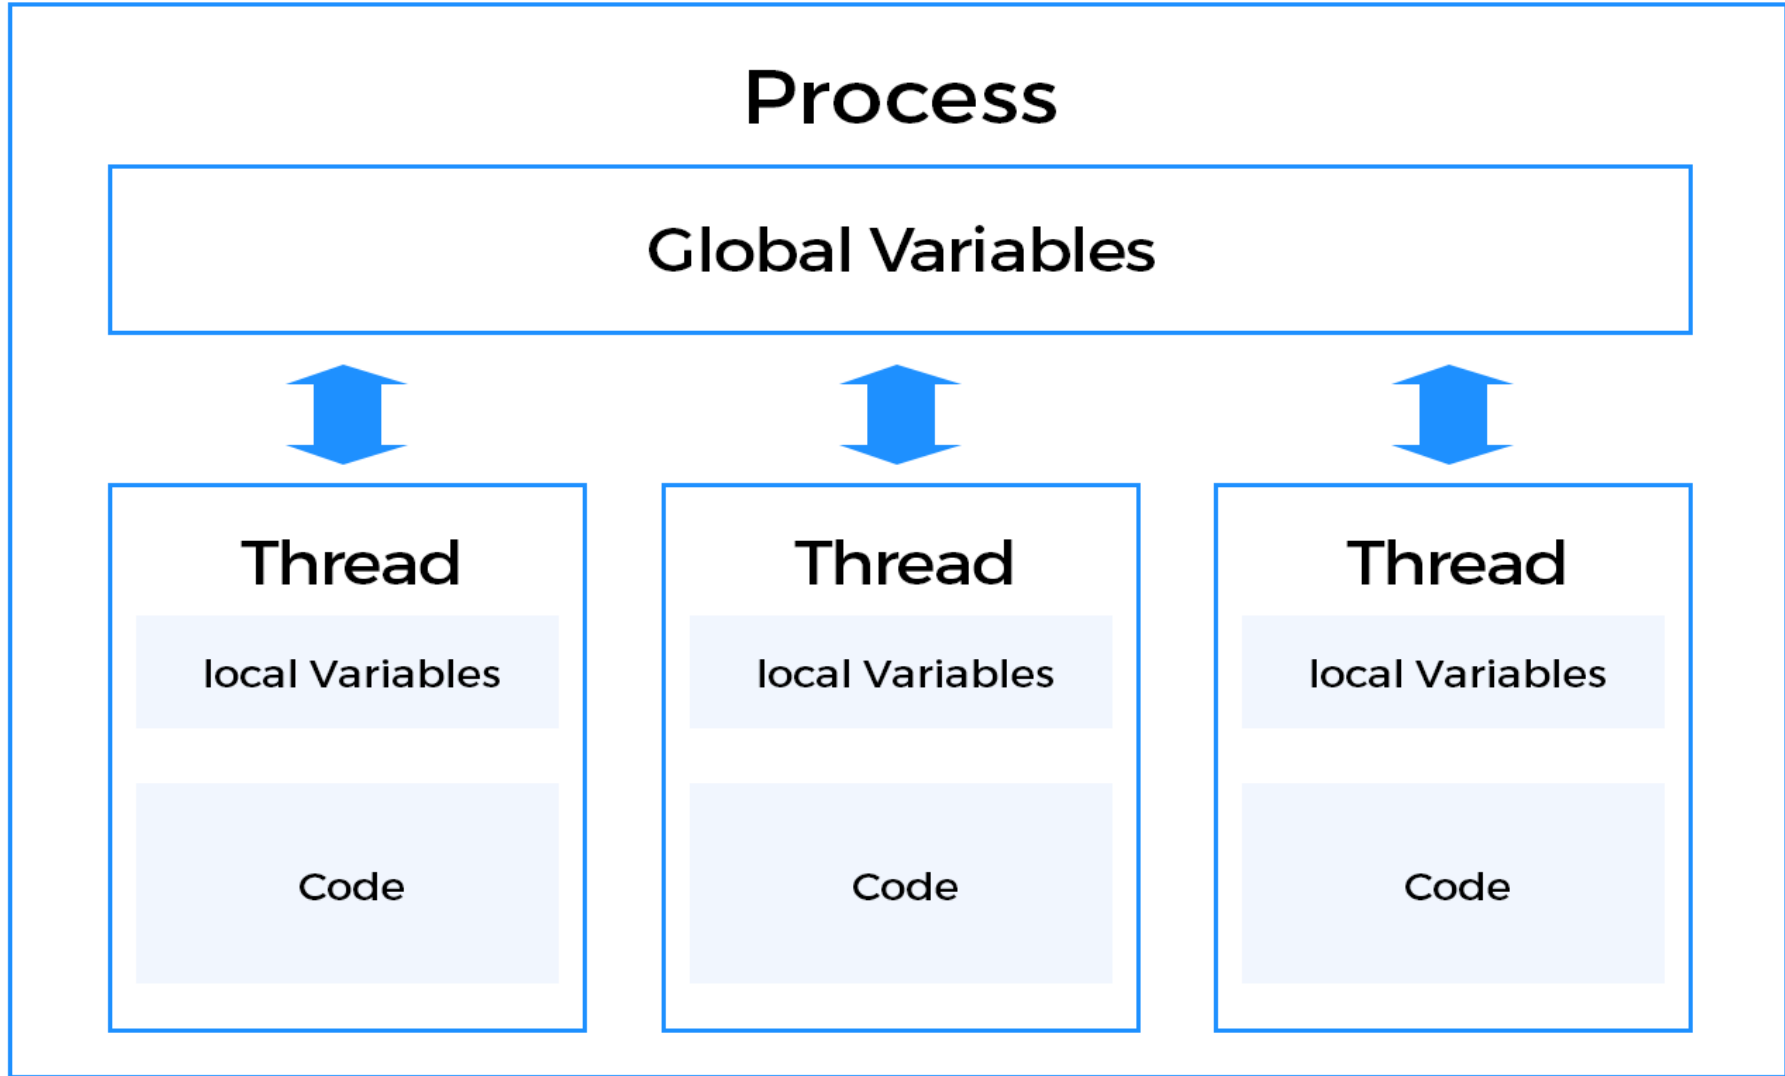


single-threaded process



multithreaded process

Process vs Thread



USER LEVEL & KERNEL LEVEL THREAD


User Level Threads	Kernel Level Thread
User level threads are faster to create and manage.	Kernel level threads are slower to create and manage.
Implementation is by a thread library at the user level.	Operating system supports creation of Kernel threads
User level thread is generic and can run on any operating system	. Kernel level thread is specific to the operating system.
Multi-threaded application cannot take advantage of multiprocessing.	Kernel routines themselves can be multithreaded.

Multithreading models are three types

Many – to – One

One – to – One

Many – to - Many



EXPLICIT THREADING - the programmer creates and manages threads.

IMPLICIT THREADING - the compilers and run-time libraries create and manage threads.

Processes

```
graph LR; A[Processes] --- B[IP Communication]; A --- C[Synchronization]; A --- D[Scheduling]; A --- E[Deadlock];
```

IP Communication

Synchronization

Scheduling

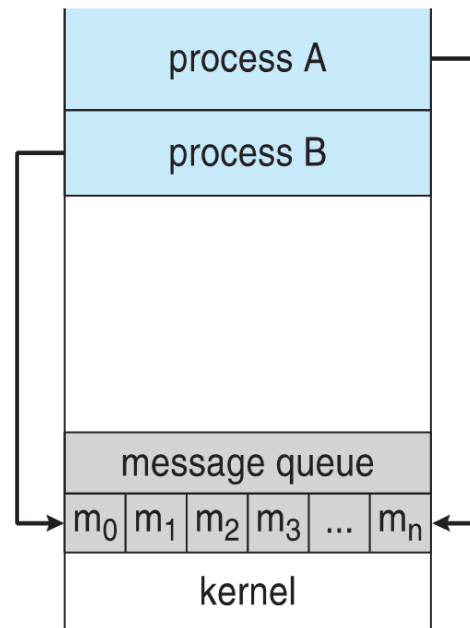
Deadlock

INTER-PROCESS COMMUNICATION


Independent Processes - neither affect other processes or be affected by other processes.

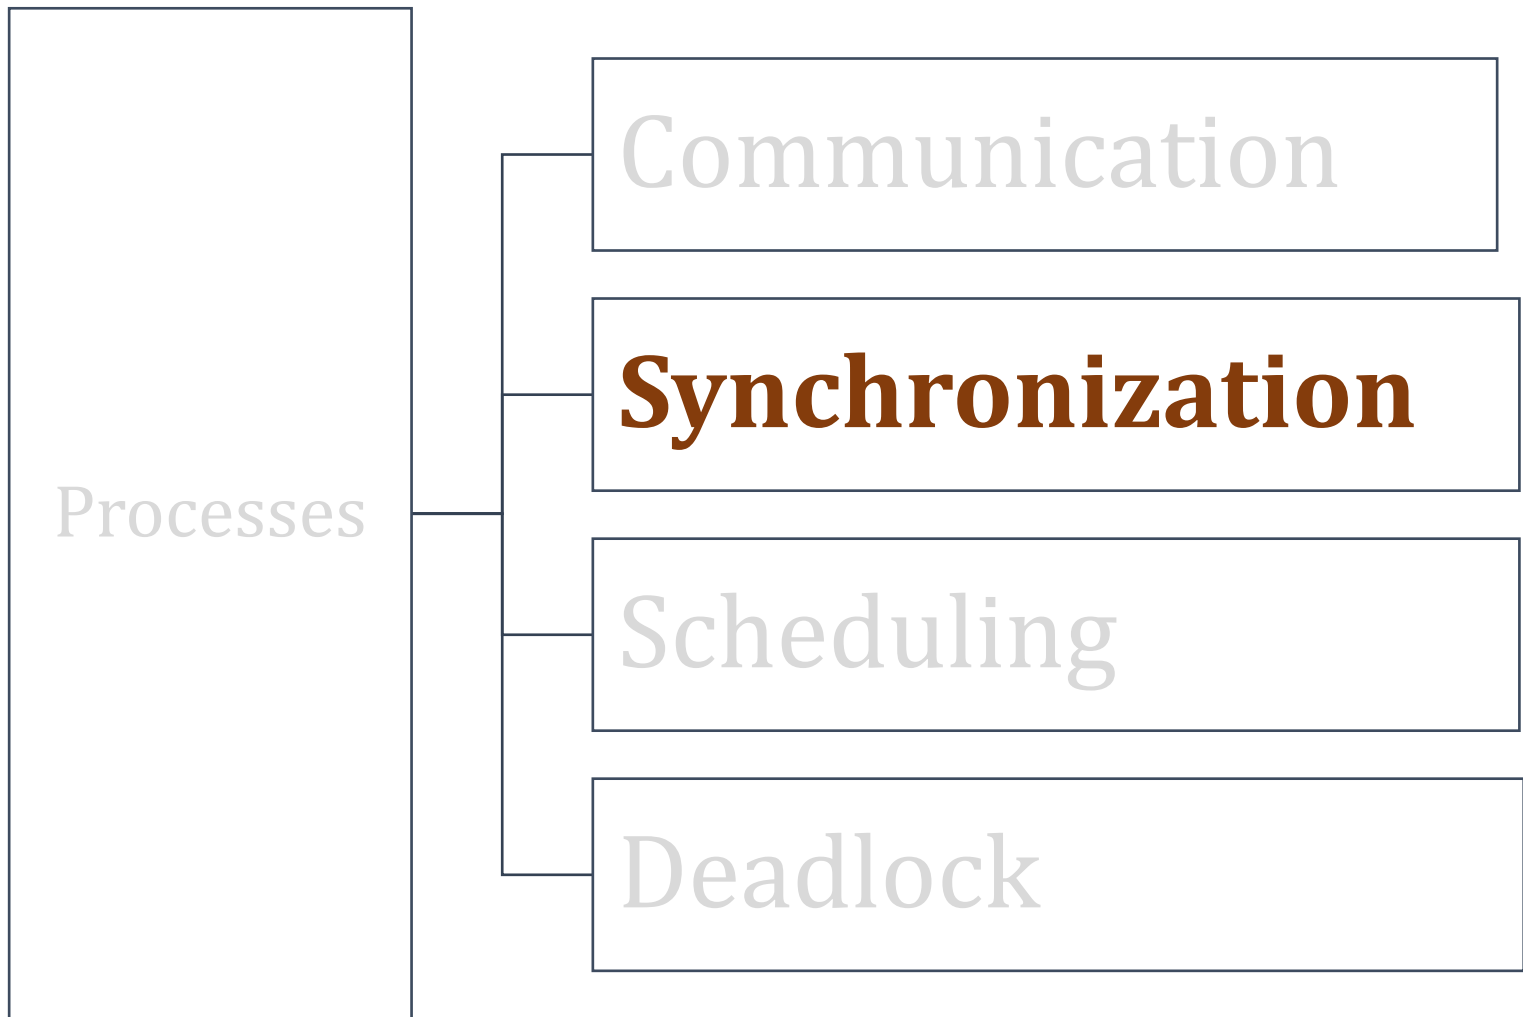
Cooperating Processes - can affect or be affected by other processes.

(a) Message passing.

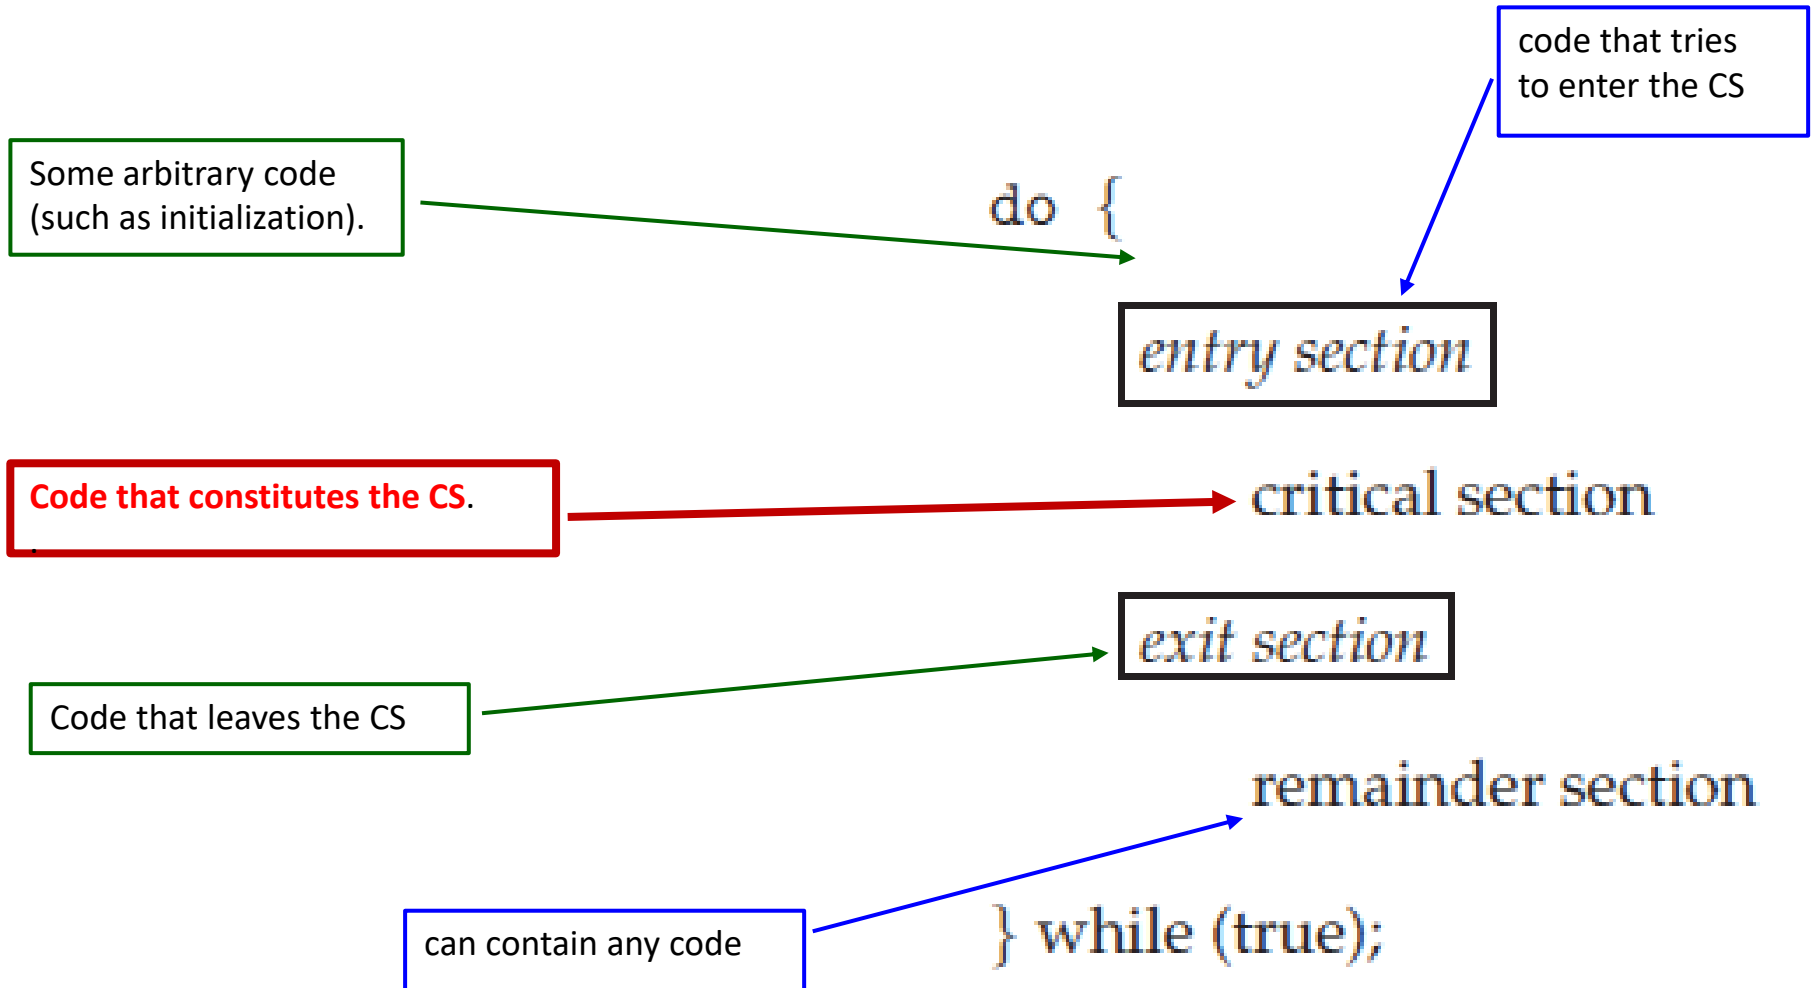


(b) Shared memory.





THE CRITICAL SECTION PROBLEM



```

int i = 7;           //global variable or shared resource
void increment() {
    i++;             // critical section of code
}
int main() {
    thread T1,T2;
    T1(increment);
    T2(increment);
}

```

What will be the final value of i after both thread have completed execution?

Case 1

Thread 1	Thread 2
get i (7)	—
increment i (7 -> 8)	—
write back i (8)	—
—	get i (8)
—	increment i (8 -> 9)
—	write back i (9)

Case 2

Thread 1	Thread 2
get i (7)	get i (7)
increment i (7 -> 8)	—
—	increment i (7 -> 8)
write back i (8)	—
—	write back i (8)

The entry- and exit-sections that surround a critical section must satisfy the following correctness requirements:


- ❑ **Mutual exclusion** - no two process/thread can be simultaneously present inside critical section at any point in time
- ❑ **Progress** - no process/thread running outside the critical section should block the other interesting process from entering into a critical section when in fact the critical section is free
- ❑ **Bounded waiting** - No process/thread should have to wait forever to enter into the critical section.

SOLUTIONS TO THE CRITICAL SECTION PROBLEM

- **Peterson's solution** - restricted to **two processes** that alternate execution between their critical sections and remainder sections.
- **Hardware solution/Synchronization Hardware** - rely on some special machine instructions
- **Mutex lock** - software to protect critical regions and avoid race conditions.
- **Semaphores** - similar to the mutex lock, provide sophisticated ways for the process to synchronize their activities.

MUTEX LOCKS / MUTUAL EXCLUSION

- a mutex is locking mechanism used to synchronize access to a resource.



- Mutex object is locked or unlocked by the process/thread requesting or releasing the resource

SEMAPHORES

There are two main types of semaphores:

- **COUNTING SEMAPHORE** – allow an arbitrary resource count. Its **value can range over an unrestricted domain**. It is used to control access to a resource that has multiple instances.
- **BINARY SEMAPHORE** – This is also known as **mutex lock**. It can have only **two values – 0 and 1**.

Its **value is initialized to 1**. It is used to implement the solution of critical section problem with multiple processes.

SEMAPHORES

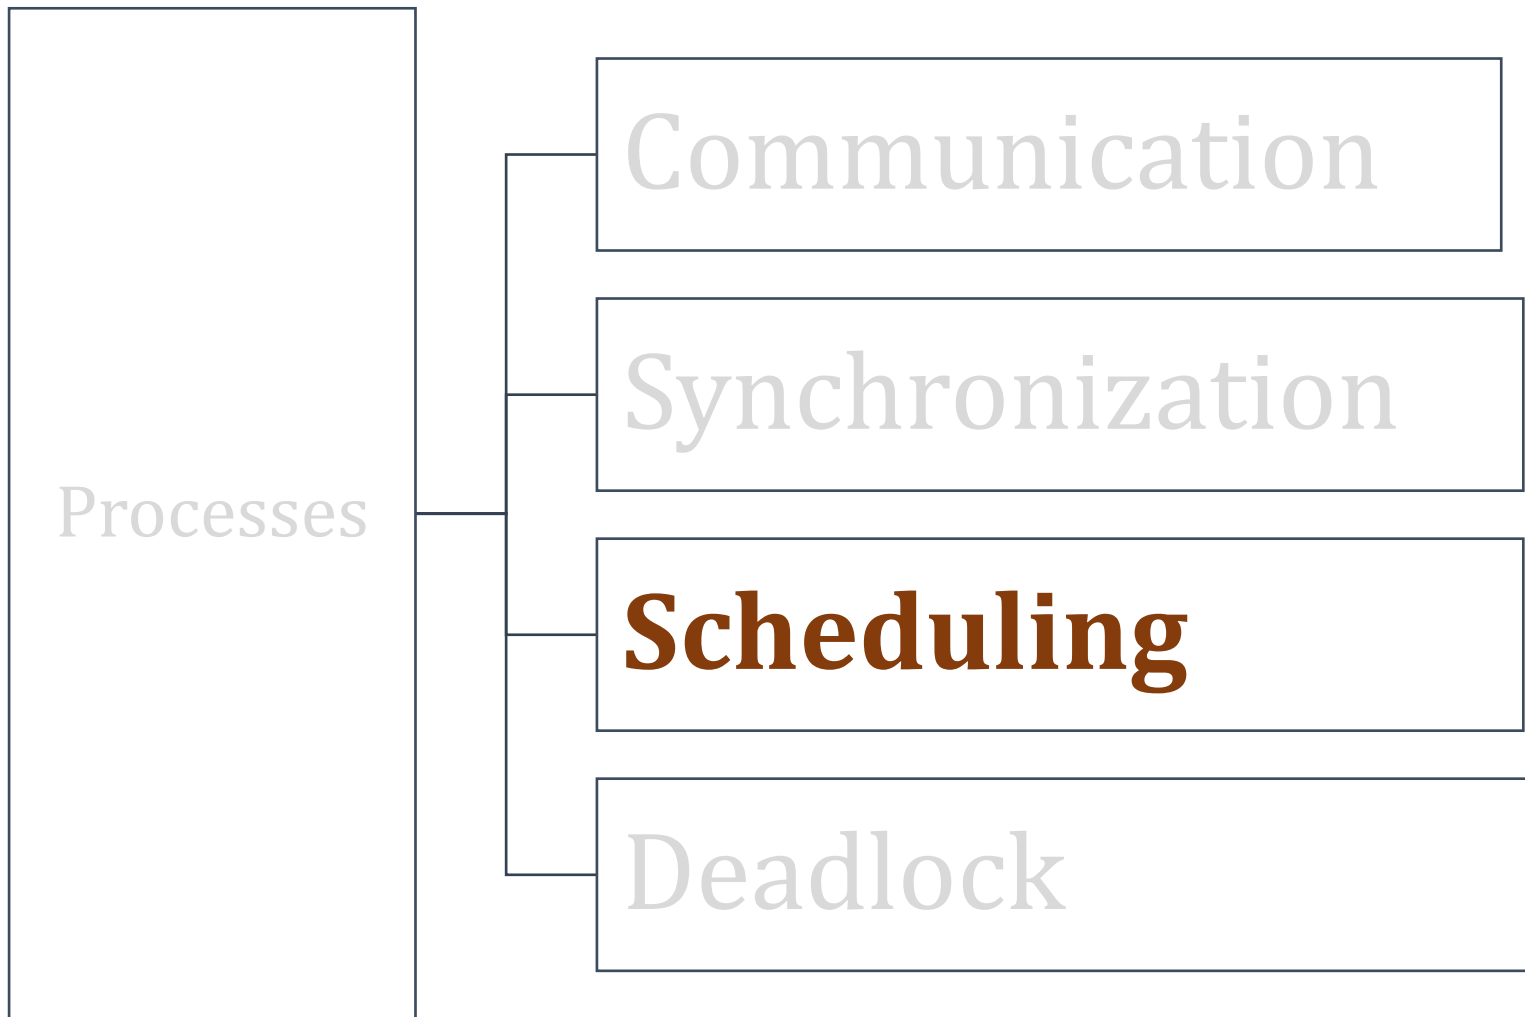
A semaphore is a signaling mechanism.

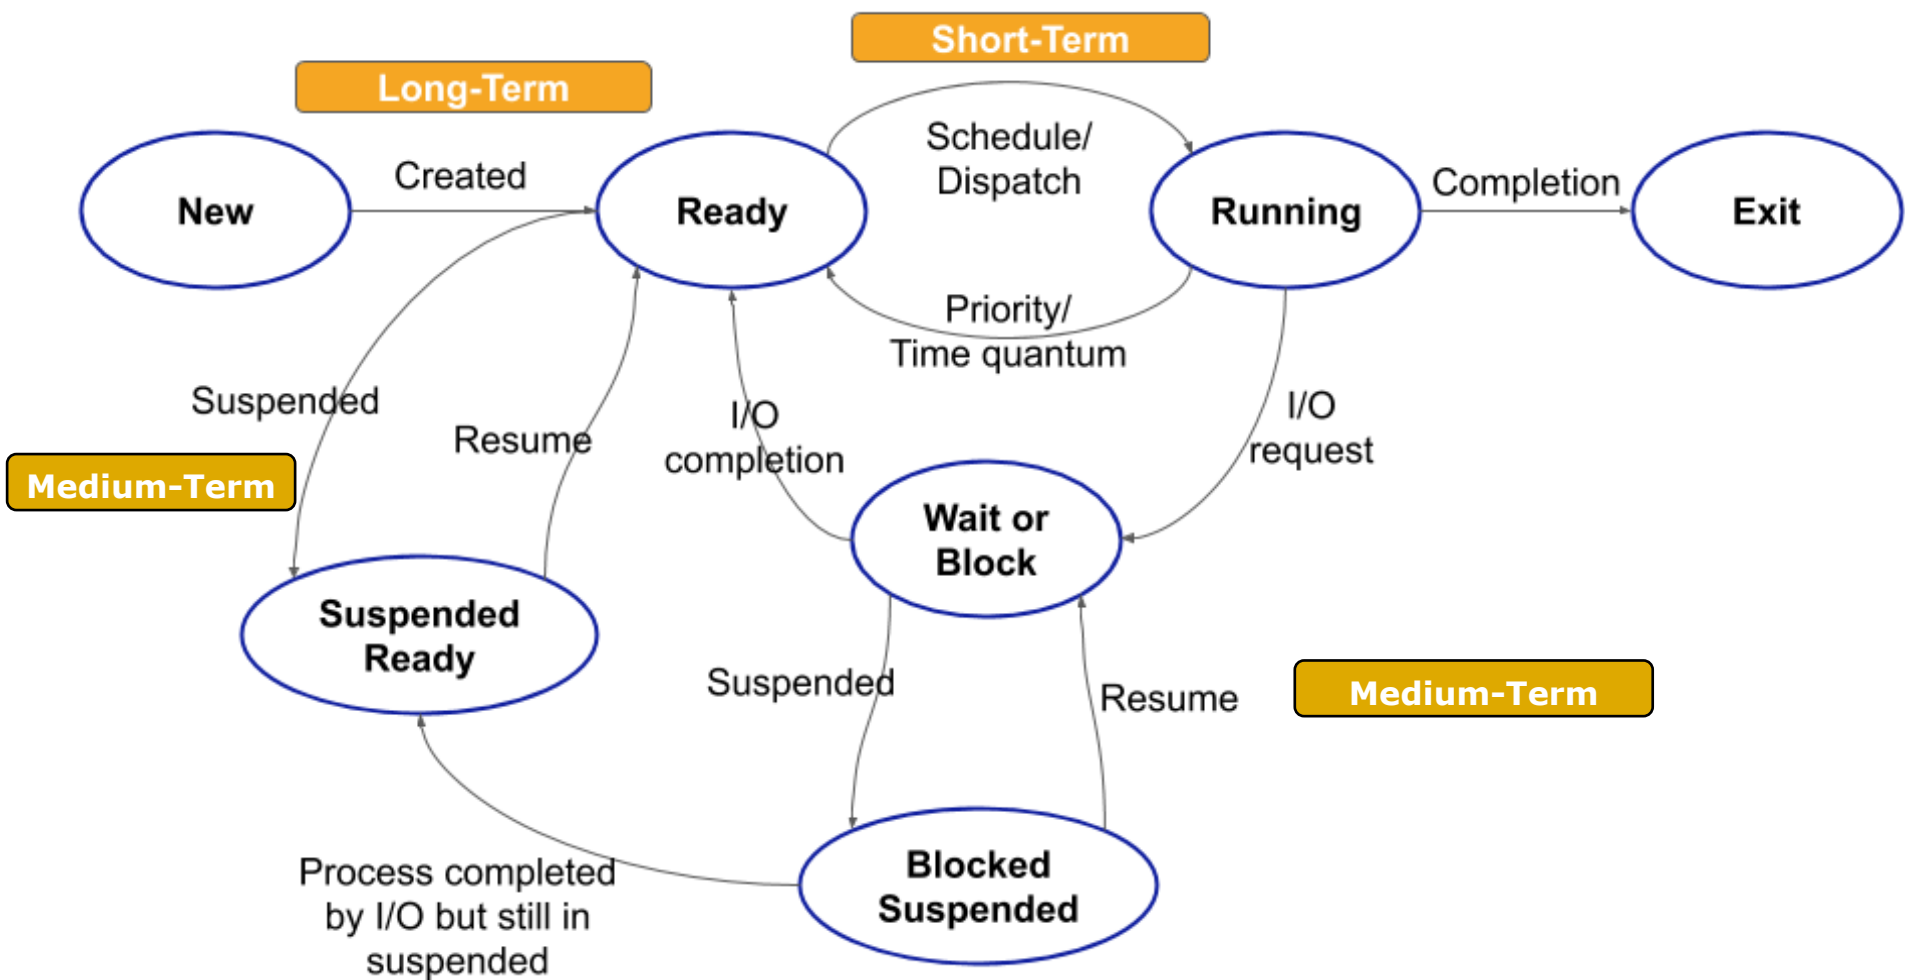
- The **wait** operation decrements the value of its argument S , if it is positive. If S is negative or zero, then no operation is performed.

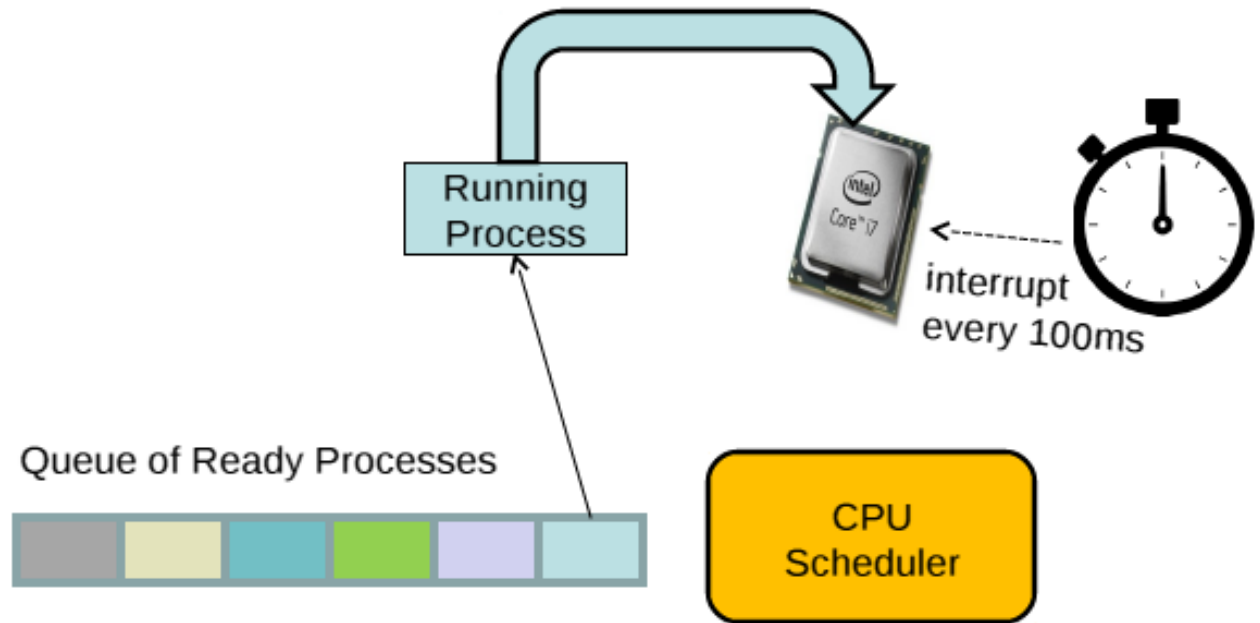
```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

- The **signal** operation increments the value of its argument S :

```
signal(S) {  
    S++;  
}
```







The scheduling in which a running process **can be interrupted** if a high priority process enters the queue and is allocated to the CPU is called **PREEMPTIVE SCHEDULING**.

The scheduling in which a running process **cannot be interrupted** by any other process is called **NON-PREEMPTIVE SCHEDULING**.

SCHEDULING ALGORITHMS

- ☐ First-Come, First-Served (FCFS) Scheduling
- ☐ Shortest-Job-First (SJF) Scheduling
- ☐ Shortest Remaining Time First
- ☐ Priority Scheduling
- ☐ Round Robin(RR) Scheduling
- ☐ Multiple-Level Queues Scheduling
- ☐ Multilevel Feedback Queue Scheduling

REAL-TIME CPU SCHEDULING

Real-time systems are those in which **the time is crucial** to their performance.

- **Priority-Based Scheduling** - each process a priority based on its importance
- **Rate-Monotonic Scheduling** - **static priority** (*the shorter the period = the higher the priority*) with preemption.
 - Missed Deadlines with Rate Monotonic Scheduling
- **Earliest-Deadline-First Scheduling** – **dynamic priority** (*the earlier the deadline = the higher the priority*) with preemption.
- **Proportional Share Scheduling** - T shares are allocated among all processes in the system.

```
graph LR; A[Processes] --- B[Communication]; A --- C[Synchronization]; A --- D[Scheduling]; A --- E[Deadlock];
```

Processes

Communication

Synchronization

Scheduling

Deadlock

CONDITIONS OF DEADLOCK

There are certain conditions that give rise to a deadlock:

- **Mutual exclusion:** Some resource types cannot allow multiple processes to access it at the same time \Rightarrow only one process at a time can use a resource
- **Hold and wait:** When all the processes are holding some resources and waiting for other resources, a deadlock may occur
- **No preemption:** If a resource cannot be pre-empted, it may lead to a deadlock.
- **Circular wait:** The Resource Allocation Graph RAG has a cycle.

Ensure that the system will *never* enter a deadlock state:

- **Deadlock prevention**
- **Deadlock avoidance**
- **Deadlock Detection**

If a system is in safe state \Rightarrow no deadlocks

If a system is in unsafe state \Rightarrow possibility of deadlock

Prevention \Rightarrow eliminating any of the four conditions

Avoidance \Rightarrow ensure that a system will never enter an unsafe state.

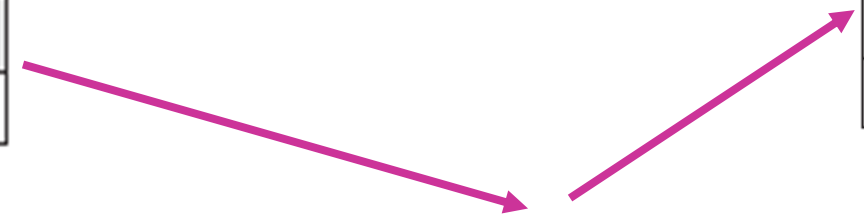
How to check if system is in safe state

Total resources

R1	R2	R3
15	8	8

Available

R1	R2	R3
3	3	2



Process	Max			Alloc			Need(Max - Alloc)		
	R1	R2	R3	R1	R2	R3	R1	R2	R3
P1	5	6	3	2	1	0	3	5	3
P2	8	5	6	3	2	3	5	3	3
P3	4	8	2	3	0	2	1	9	0
P4	7	4	3	3	2	0	4	2	3
P5	4	3	3	1	0	1	3	3	2

Available = Available - Request[i];

Allocation[i] = Allocation[i] + Request[i];

Need[i] = Need[i] - Request[i];

MEMORY

Objectives of a Memory Management (MM) System

- Protection
- Relocation
- Sharing
- Logical Organization of memory
- Physical Organization of memory

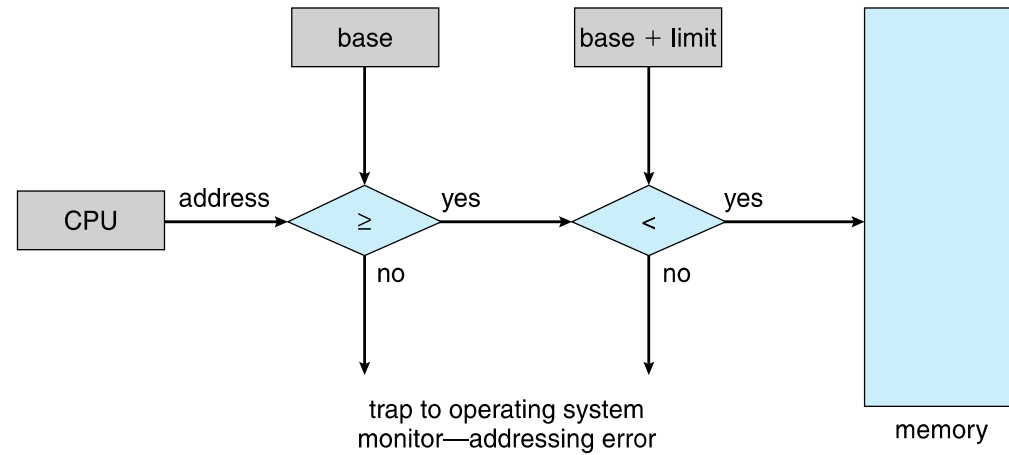
Logical address – generated by the CPU; also referred to as *virtual address*.

When the process starts executing, *relative or logical addresses* are generated by the CPU.


Physical address – address seen by the memory unit (*the absolute addresses*)

A pair of **base register / relocation register** and **limit registers** define the *logical address space*


- To translate all the memory references of the process, **there must be an origin or base address in the memory.**
- All relative addresses generated are added in this base address to get the new location in the memory.



Swapping



Relocation



Logical Organization of Memory: Allocation


The memory allocation can be classified into two methods:

- **contiguous memory allocation** - assigns consecutive memory blocks to a process.
- **non-contiguous memory allocation** - assigns different blocks of memory in a nonconsecutive manner to a process.

CONTIGUOUS ALLOCATION

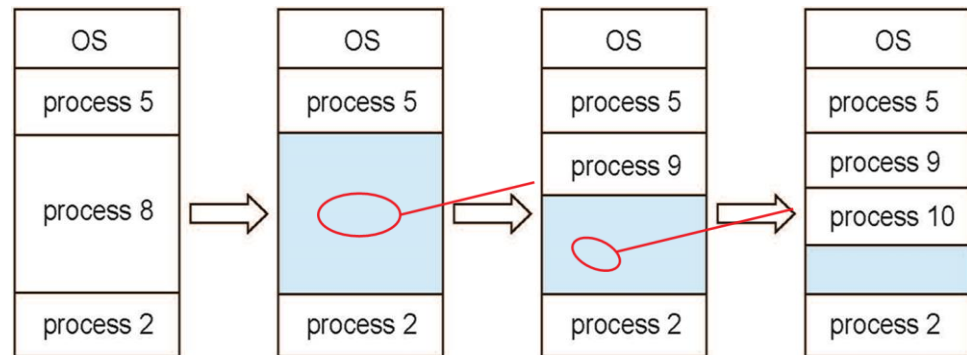
Fixed/Static Partitioning

- fixed partitions can be of *equal* or *unequal* sizes.



Variable/Dynamic Partitioning

- list of free memory blocks (**holes**) to find a hole of a suitable size whenever a process needs to be loaded into memory.



Placement Algorithm:

First-fit: Allocate the first hole that's big enough

Best-fit: Allocate the smallest hole that's big enough

Worst-fit: Allocate the largest hole


Fragmentation

Internal fragmentation

- occurs in fixed size blocks, because the last allocated process is not completely filled.


Solution:

- can be reduced by using variable sized memory blocks rather than fixed sized.



External fragmentation

- occurs with variable size segments, because some holes in memory will be too small to use.



Solutions:

Compaction - moving all occupied areas of storage to one end of memory. This leaves one big hole.

Non-contiguous memory allocation:
Segmentation and **Paging**.

To Be Continued